

# Introduction To Data Analysis Using Pandas on Watson Studio

Kunal Malhotra

kunal.malhotra1@ibm.com

@kunal\_forever



# Agenda

- What is Data Analysis ?
- Introduction to Pandas.
- Advantages of Pandas.
- Introduction to Watson Studio.
- How to perform data analysis with pandas.

# Data Analysis

- What is it?
  - Apply logical techniques to describe, condense, recap and evaluate Data and illustrate information.
- Goal of Data Analysis:
  - Discover useful information.
  - Provide insights.
  - Suggest conclusions.
  - Support decision Making.

# Introduction To Pandas.

- **Pandas** is a python package for **data analysis**.
- It Provides built-in data structures which simplify the manipulation and analysis of data set.
- Pandas is easy to use and powerful, but “with great power comes great responsibility”.
- <http://pandas.pydata.org/pandas-docs/stable/>

# Pandas: Essential Concepts

- **Series**
- **DataFrame**
- **Reading and Writing Files**
- **Aggregating and Grouping in Pandas**
- **Time Series analysis using pandas**
- **Visualization in pandas**

# Introduction To Series In Pandas

**Series** in pandas is an object which is similar to python built-in list data structure, but differs from it because it has associated label with each **element** or better know as **index**.

```
In [1]:
```

```
import pandas as pd
```

```
In [2]:
```

```
series = pd.Series([1,2,3,4,5,6,7,8,9])
```

```
In [3]:
```

```
series
```

```
Out[3]:
```

```
0    1
1    2
2    3
3    4
4    5
5    6
6    7
7    8
8    9
dtype: int64
```

# Understanding Series

- **Index** is leftward and **Values** are to the right. (Note:- If **Index** is not provided explicitly, then pandas creates **RangeIndex** starting from **0 to N - 1**, where N is the total number of elements.)
- Each series object has a data type (**dtype**), in the example on the right data type is **int64**.

In [1]:

```
import pandas as pd
```

In [2]:

```
series = pd.Series([1,2,3,4,5,6,7,8,9])
```

In [3]:

```
series
```

Out[3]:

0	1
1	2
2	3
3	4
4	5
5	6
6	7
7	8
8	9

dtype: int64

# Operation On Pandas Series

- Pandas series have attributes to extract its **values** and **labels**.
- Elements can be retrieved by their **labels(index)**.

In [4]:

```
series.index
```

Out[4]:

```
RangeIndex(start=0, stop=9, step=1)
```

In [5]:

```
series.values
```

Out[5]:

```
array([1, 2, 3, 4, 5, 6, 7, 8, 9])
```

In [6]:

```
series[4]
```

Out[6]:

```
5
```



# Operation On Pandas Series Continued ..

- **Labels (index)** can be provided explicitly.
- Elements can be retrieved by the provided **labels(index)**.

```
In [1]: import pandas as pd
```

```
In [7]: series = pd.Series([1,2,3,4,5,6,7,8,9], index=['a','b','c','d','e','f','g','h','i'])
```

```
In [8]: series
```

```
Out[8]: a    1  
       b    2  
       c    3  
       d    4  
       e    5  
       f    6  
       g    7  
       h    8  
       i    9  
       dtype: int64
```

```
In [9]: series['f']
```

```
Out[9]: 6
```

# Operation On Pandas Series Continued ..

- It is easy to retrieve several elements by their indexes or make group assignment.

```
In [1]: import pandas as pd
```

```
In [7]: series = pd.Series([1,2,3,4,5,6,7,8,9], index=['a','b','c','d','e','f','g','h','i'])
```

```
In [10]: series[['a','b','f']]
```

```
Out[10]: a    1  
         b    2  
         f    6  
         dtype: int64
```

```
In [11]: series[['a','b','c']] = 0
```

```
In [12]: series
```

```
Out[12]: a    0  
         b    0  
         c    0  
         d    4  
         e    5  
         f    6  
         g    7  
         h    8  
         i    9  
         dtype: int64
```

# Operation On Pandas Series Continued ..

- Filtering and Maths operations are easy as well.

```
In [1]: import pandas as pd

In [13]: series = pd.Series([1,2,3,4,5,6,7,8,9], index=['a','b','c','d','e','f','g','h','i'])

In [14]: series[series > 5]

Out[14]: f    6
         g    7
         h    8
         i    9
         dtype: int64

In [15]: series[series > 5] * 2

Out[15]: f    12
         g    14
         h    16
         i    18
         dtype: int64
```

# Operation On Pandas Series Continued ..

- Series is very similar to dictionary, where key is an index and value is an element. Hence, we have generated a pandas series with python dictionary.

```
In [1]: import pandas as pd

In [16]: series = pd.Series({'a':5,'b':6,'c':7,'d':8})

In [17]: series
Out[17]: a    5
         b    6
         c    7
         d    8
         dtype: int64

In [18]: 'd' in series
Out[18]: True
```

# Introduction To DataFrame in Pandas

- DataFrame in pandas is a two-dimensional data structure, i.e., data is aligned in a tabular fashion. It has rows and columns. Each column in a DataFrame is a series object, rows consist of elements in Series.
- DataFrame can be constructed using built-in Python dictionary.

# DataFrame in Pandas

```
In [1]: import pandas as pd
```

```
In [19]: data_frame = pd.DataFrame({  
    'country': ['USA', 'UK', 'UAE', 'Germany', 'Russia'],  
    'population (Million)': ['325.7', '65.64', '9.27', '82.67', '144.3'],  
    'Square Area (km²)': ['98340000', '242,495', '83,600', '357,376', '17100000']  
})
```

```
In [20]: data_frame
```

```
Out[20]:
```

	Square Area (km²)	country	population (Million)
0	98340000	USA	325.7
1	242,495	UK	65.64
2	83,600	UAE	9.27
3	357,376	Germany	82.67
4	17100000	Russia	144.3

```
In [24]: data_frame['country']
```

```
Out[24]:
```

```
0    USA  
1    UK  
2    UAE  
3  Germany  
4   Russia  
Name: country, dtype: object
```

```
In [25]: type(data_frame['country'])
```

```
Out[25]: pandas.core.series.Series
```

# Understanding DataFrame

- **DataFrame** object has **2 indexes**: column index and row index.

```
In [1]: import pandas as pd

In [19]: data_frame = pd.DataFrame({
    'country': ['USA', 'UK', 'UAE', 'Germany', 'Russia'],
    'population (Million)': ['325.7', '65.64', '9.27', '82.67', '144.3'],
    'Square Area (km²)': ['98340000', '242,495', '83,600', '357,376', '17100000']
})

In [27]: data_frame.columns
Out[27]: Index(['Square Area (km²)', 'country', 'population (Million)'], dtype='object')

In [29]: data_frame.index
Out[29]: RangeIndex(start=0, stop=5, step=1)
```

In the above image the DataFrame has 5 elements from 0 to 4.

# Operation On Pandas DataFrame

- There are numerous ways to provide row index explicitly, for example you can provide index when creating a DataFrame or do it “on the fly” during runtime.

```
In [2]: import pandas as pd
```

```
In [3]: data_frame = pd.DataFrame({
    'country': ['USA', 'UK', 'UAE', 'Germany', 'Russia'],
    'population (Million)': ['325.7', '65.64', '9.27', '82.67', '144.3'],
    'Square Area (km²)': ['98340000', '242,495', '83,600', '357,376', '17100000']
}, index=['First', 'Second', 'Third', 'Fourth', 'Fifth'])
```

```
In [4]: data_frame
```

```
Out[4]:
```

	Square Area (km²)	country	population (Million)
First	98340000	USA	325.7
Second	242,495	UK	65.64
Third	83,600	UAE	9.27
Fourth	357,376	Germany	82.67
Fifth	17100000	Russia	144.3



# Operation On Pandas DataFrame Continued ..

- Row access using index can be performed in several ways. The two most important ones are:-
  - using **.loc** and providing index label
  - using **.iloc** and providing index number

```
In [2]: import pandas as pd
```

```
In [3]: data_frame = pd.DataFrame({
        'country': ['USA', 'UK', 'UAE', 'Germany', 'Russia'],
        'population (Million)': ['325.7', '65.64', '9.27', '82.67', '144.3'],
        'Square Area (km²)': ['98340000', '242,495', '83,600', '357,376', '17100000']
    }, index=['First', 'Second', 'Third', 'Fourth', 'Fifth'])
```

```
In [5]: data_frame.loc['Fourth']
```

```
Out[5]: Square Area (km²)    357,376
country                Germany
population (Million)    82.67
Name: Fourth, dtype: object
```

```
In [11]: data_frame.iloc[3]
```

```
Out[11]: Square Area (km²)    357,376
country                Germany
population (Million)    82.67
Name: Fourth, dtype: object
```

# Operation On Pandas DataFrame Continued ..

- Selection of particular rows and columns

```
In [13]: data_frame.loc[['Fourth','Fifth'], 'population (Million)']
```

```
Out[13]: Fourth      82.67  
        Fifth      144.3  
        Name: population (Million), dtype: object
```

- **.loc** takes **2 arguments**: index list and column list, **slicing** operation is supported

```
In [16]: data_frame.loc[['Fourth','Fifth'], :]
```

```
Out[16]:
```

	Square Area (km²)	country	population (Million)
Fourth	357,376	Germany	82.67
Fifth	17100000	Russia	144.3

# Operation On Pandas DataFrame Continued ..

- Filtering can be performed using Boolean arrays.

```
In [26]: import pandas as pd
```

```
In [52]: data_frame = pd.DataFrame({ 'country':['USA','UK','UAE','Germany','Russia'],  
                                     'population':[325.7,65.64,9.27,82.67,144.3],  
                                     'square':[98340000,242495,83600,357376,17100000]},  
                                    index=['First','Second','Third','Fourth','Fifth'])
```

```
In [58]: data_frame[data_frame['population'] > 100][['country','square']]
```

```
Out[58]:
```

	country	square
First	USA	98340000
Fifth	Russia	17100000

# Operation On Pandas DataFrame Continued ..

- Adding a new column, for example adding population density column.

```
In [26]: import pandas as pd
```

```
In [52]: data_frame = pd.DataFrame({ 'country': ['USA', 'UK', 'UAE', 'Germany', 'Russia'],  
                                     'population': [325.7, 65.64, 9.27, 82.67, 144.3],  
                                     'square': [98340000, 242495, 83600, 357376, 17100000],  
                                     index= ['First', 'Second', 'Third', 'Fourth', 'Fifth']})
```

```
In [55]: data_frame['density'] = data_frame['population'] / data_frame['square'] * 1000000
```

```
In [56]: data_frame
```

```
Out[56]:
```

	country	population	square	density
First	USA	325.70	98340000	3.311979
Second	UK	65.64	242495	270.685994
Third	UAE	9.27	83600	110.885167
Fourth	Germany	82.67	357376	231.324991
Fifth	Russia	144.30	17100000	8.438596

# Operation On Pandas DataFrame Continued ..

- Deleting a column.

```
In [26]: import pandas as pd
```

```
In [52]: data_frame = pd.DataFrame({ 'country':['USA','UK','UAE','Germany','Russia'],  
                                     'population':[325.7,65.64,9.27,82.67,144.3],  
                                     'square':[98340000,242495,83600,357376,171000000]},  
                                   index=['First','Second','Third','Fourth','Fifth'])
```

```
In [59]: data_frame.drop(['density'], axis='columns')
```

Out[59]:

	country	population	square
First	USA	325.70	98340000
Second	UK	65.64	242495
Third	UAE	9.27	83600
Fourth	Germany	82.67	357376
Fifth	Russia	144.30	171000000

# Reading and Writing Files in Pandas

- Pandas support many popular file formats including **CSV, XML, HTML, Excel, SQL, JSON** many more.
- For example, writing a dataframe to a **CSV file** and then reading it.

```
In [2]: import pandas as pd
```

```
In [3]: data_frame = pd.DataFrame({ 'country': ['USA', 'UK', 'UAE', 'Germany', 'Russia'],  
                                   'population': [325.7, 65.64, 9.27, 82.67, 144.3],  
                                   'square': [98340000, 242495, 83600, 357376, 171000000]},  
                                   index= ['First', 'Second', 'Third', 'Fourth', 'Fifth'])
```

```
In [4]: data_frame.to_csv('filename.csv')
```

```
In [5]: data_frame2 = pd.read_csv('filename.csv', sep=',')
```

```
In [6]: data_frame2
```

```
Out[6]:
```

	Unnamed: 0	country	population	square
0	First	USA	325.70	98340000
1	Second	UK	65.64	242495
2	Third	UAE	9.27	83600
3	Fourth	Germany	82.67	357376
4	Fifth	Russia	144.30	171000000

## Reading And Writing Files in Pandas Continued ..

- **`to_csv`** method takes many arguments for example, **separator character**.
- As shown on the previous slides, named argument `sep` points to a separator character in CSV file called *filename.csv*.
- There are many different ways to construct DataFrame from external sources, for example using `read_sql` method pandas can perform SQL query and store results inside a new DataFrame instance

# Aggregating and Grouping in Pandas

- Grouping is probably one of the most popular methods in data analysis. If you want to group data in pandas you have to use **.groupby** method.
- In order to demonstrate aggregates and grouping in pandas I decided to choose popular Titanic dataset (<https://yadi.sk/d/TfhJdE2k3EyALt>)



## Loading Data

```
In [8]: import sys
import types
import pandas as pd
from botocore.client import Config
import ibm_boto3
def __iter__(self): return 0
# @hidden_cell
# The following code accesses a file in your IBM Cloud Object Storage. It includes your credentials.
# You might want to remove those credentials before you share your notebook.
client_64dd1b2268ad43e4b4905b0a0968e5f6 = ibm_boto3.client(service_name='s3',
    ibm_api_key_id='S1P54okzAoUahk3AomMOYgAoFC-bQRNQjR99P8WdnUGf',
    ibm_auth_endpoint="https://iam.ng.bluemix.net/oidc/token",
    config=Config(signature_version='oauth'),
    endpoint_url='https://s3-api.us-geo.objectstorage.service.networklayer.com')

body = client_64dd1b2268ad43e4b4905b0a0968e5f6.get_object(Bucket='dataanalysiswithpandas-donotdelete-pr-yqtectuhrngnyckb',Key='titanic.csv')['Body']
# add missing __iter__ method, so pandas accepts body as file-like object
if not hasattr(body, "__iter__"): body.__iter__ = types.MethodType( __iter__, body )

df_data_1 = pd.read_csv(body)
df_data_1.head()
```

Out[8]:

	PassengerID	Name	PClass	Age	Sex	Survived	SexCode
0	1	Allen, Miss Elisabeth Walton	1st	29.00	female	1	1
1	2	Allison, Miss Helen Loraine	1st	2.00	female	0	1
2	3	Allison, Mr Hudson Joshua Creighton	1st	30.00	male	0	0
3	4	Allison, Mrs Hudson JC (Bessie Waldo Daniels)	1st	25.00	female	0	1
4	5	Allison, Master Hudson Trevor	1st	0.92	male	1	0



# Aggregating and Grouping in Pandas Continued ..

- Let's calculate how many passengers (women and men) survived and how many did not, we will use **.groupby** as stated above.

```
In [9]: print(df_data_1.groupby(['Sex', 'Survived'])['PassengerID'].count())
```

```
Sex    Survived
female 0         154
        1         308
male   0         709
        1         142
Name: PassengerID, dtype: int64
```

- Now, let's analyze the same data by cabin class

```
In [10]: print(df_data_1.groupby(['PClass', 'Survived'])['PassengerID'].count())
```

```
PClass Survived
*      0         1
1st    0        129
        1        193
2nd    0        160
        1        119
3rd    0        573
        1        138
Name: PassengerID, dtype: int64
```



# Time series analysis using Pandas

Pandas was created to analyze time series data. In order to illustrate how easy it is, We will use Apple's last 5 year stock prices ([https://yadi.sk/d/po\\_usmXT3ExwzV](https://yadi.sk/d/po_usmXT3ExwzV)).

```
In [26]: import sys
import types
import pandas as pd
from botocore.client import Config
import ibm_boto3

def __iter__(self): return 0

# @hidden_cell
# The following code accesses a file in your IBM Cloud Object Storage. It includes your credentials.
# You might want to remove those credentials before you share your notebook.
client_64ddb2268ad43e4b4905b0a0968e5f6 = ibm_boto3.client(service_name='s3',
    ibm_api_key_id='S1P54okzAoUahk3AomMOYgAoFC-bQRNQjR99P8WdnUGf',
    ibm_auth_endpoint="https://iam.ng.bluemix.net/oidc/token",
    config=Config(signature_version='oauth'),
    endpoint_url='https://s3-api.us-geo.objectstorage.service.networklayer.com')

body = client_64ddb2268ad43e4b4905b0a0968e5f6.get_object(Bucket='dataanalysiswithpandas-donotdelete-pr-yqtectuhrngnyckb',Key='apple.csv')['Body']
# add missing __iter__ method, so pandas accepts body as file-like object
if not hasattr(body, "__iter__"): body.__iter__ = types.MethodType( __iter__, body )

df_data_1 = pd.read_csv(body, index_col='Date', parse_dates=True)
df_data_1.head()
```

Out[26]:

	Open	High	Low	Close	Volume	Adj Close
Date						
2017-02-22	136.429993	137.119995	136.110001	137.110001	20745300	137.110001
2017-02-21	136.229996	136.750000	135.979996	136.699997	24265100	136.699997
2017-02-17	135.100006	135.830002	135.100006	135.720001	22084500	135.720001
2017-02-16	135.669998	135.899994	134.839996	135.350006	22118000	135.350006
2017-02-15	135.520004	136.270004	134.619995	135.509995	35501600	135.509995



# Time series analysis using Pandas Continued ..

The image on the previous slide shows a sorted DataFrame with *DatetimeIndex* by *Date* column. If datetime column is different from ISO8601 format, then you have to use built-in pandas function ***pandas.to\_datetime*** to format it.

- Let's now calculate mean closing price.

```
In [30]: df_data_1.loc['2012-Feb', 'Close'].mean()  
Out[30]: 528.4820021999999
```

- But what about specific time period?

```
In [31]: df_data_1.loc['2012-Feb': '2015-Feb', 'Close'].mean()  
Out[31]: 430.43968317018414
```

# Time series analysis using Pandas Continued ..

- Let's calculate mean of closing price by weeks.

```
In [32]: df_data_1.resample('W')['Close'].mean()
```

```
Out[32]: Date
2012-02-26    519.399979
2012-03-04    538.652008
2012-03-11    536.254004
2012-03-18    576.161993
2012-03-25    600.990001
2012-04-01    609.698003
2012-04-08    626.484993
2012-04-15    623.773999
2012-04-22    591.718002
2012-04-29    590.536005
2012-05-06    579.831995
2012-05-13    568.814001
2012-05-20    543.593996
2012-05-27    563.283995
2012-06-03    572.539994
2012-06-10    570.124002
2012-06-17    573.029991
2012-06-24    583.739993
2012-07-01    574.070004
2012-07-08    601.937489
2012-07-15    606.080008
2012-07-22    607.746011
2012-07-29    587.951999
2012-08-05    607.217999
2012-08-12    621.150003
2012-08-19    635.394003
2012-08-26    663.185999
2012-09-02    670.611995
2012-09-09    675.477503
2012-09-16    673.476007
...
2016-08-07    105.934003
2016-08-14    108.258000
2016-08-21    109.304001
```



# Time series analysis using Pandas Continued ..

- Resampling is a very powerful tool when it comes to time series analysis.
- Resampling can be defined as a number of string aliases, given to useful common time series frequencies. In the above image, I am using “W”.
- For more information on resampling, refer pandas official documentation  
(<http://pandas.pydata.org/pandas-docs/stable/timeseries.html#offset-aliases>)


# Visualization in Pandas

For visualization pandas use library called **matplotlib**.

Let's see how Apple stock prices change over time on a graph:

- Taking Closing price between Feb, 2012 and Feb, 2017

```
In [35]: import matplotlib.pyplot as plt
df_data_11 = df_data_1.loc['2012-Feb':'2017-Feb', ['Close']]
df_data_11.plot()
plt.show()
```



## Visualization in Pandas Continued ..

Values of X-axis are represented by index values of DataFrame (by default if not provide explicitly), Y-axis represents the closing price.




# Advantage Of Pandas

Pandas have several advantages over other solutions such as NumPy, Statsmodels, SciPy etc.

1. A Pandas DataFrame can have non-homogeneous data i.e you can have different data types(int, float, string, datetime etc.) all in one place.
2. Good IO capabilities.
3. Pandas have built in functionality for a lot of common data-processing applications, for example:- easy group by syntax, easy joins (which is extremely efficient in pandas), rolling windows.

# Introduction To IBM Watson Studio


 IBM Watson

ProjectsToolsCommunityServices

DocsSupportManage

KM

Get started ▾





Welcome Kunal!


Watson Studio is part of IBM Watson.


Try out other IBM Watson apps.


Get started with key tasks


  
New project

  
Refine data

  
New notebook









  
Deep learning


  
New Modeler flow

  
New model

Hide ▾

Recently updated projects [View all \(6\)](#) [+ New project](#)

NAME	ROLE	COLLABORATORS	DATE CREATED	LAST UPDATED
Okay	Admin		Apr 09, 2018	Apr 10, 2018
BreakIntoAI	Admin		Apr 09, 2018	Apr 09, 2018
LabsKM	Admin	   +1	Mar 14, 2018	Apr 04, 2018
Data Analysis Using Pandas	Admin		Mar 24, 2018	Mar 24, 2018
JokeSentimentAnalysis	Admin	 	Mar 12, 2018	Mar 12, 2018



# IBM Watson Studio

- Watson Studio is a cloud based development and deployment environment for Machine Learning, Deep Learning, Data Governance and Data Exploration.
- A platform build for business analyst, data engineer, data scientist and developer to simplify their tasks with an intuitive UI and provide massive computing power.
- A platform where insights can be traced back to models, projects, notebooks and data sources and where model can evolve and automatically update themselves.

# Advantages Of IBM Watson Studio For Data Analysis.

IBM Watson Studio is an IDE (Integrated Development Environment), available on IBM Cloud, with many advantages such as:-

1. Offering easier access to large amounts of data while decreasing the total time of analysis.
2. Rapid development experience with access to tools and utilities that break down language barriers.
3. Ability to integrate and connect to multiple data sources, allows refining, and accessing big data engines.

May the force of  
Pandas Be With  
You :)

IBM  
**CODE**

